

The First Few Tokens Are All You Need: An Efficient and Effective *Unsupervised Prefix* Fine-Tuning Method for Reasoning Models

Ke Ji^{*,1,2}, Jiahao Xu^{*,1}, Tian Liang^{*,1}, Qiuzhi Liu^{*,1}, Zhiwei He¹, Xingyu Chen¹,
Xiaoyuan Liu¹, Zhijie Wang¹, Junying Chen², Benyou Wang^{†,2},
Zhaopeng Tu^{†,1}, Haitao Mi¹, and Dong Yu¹

¹Tencent AI Lab

²The Chinese University of Hong Kong, Shenzhen

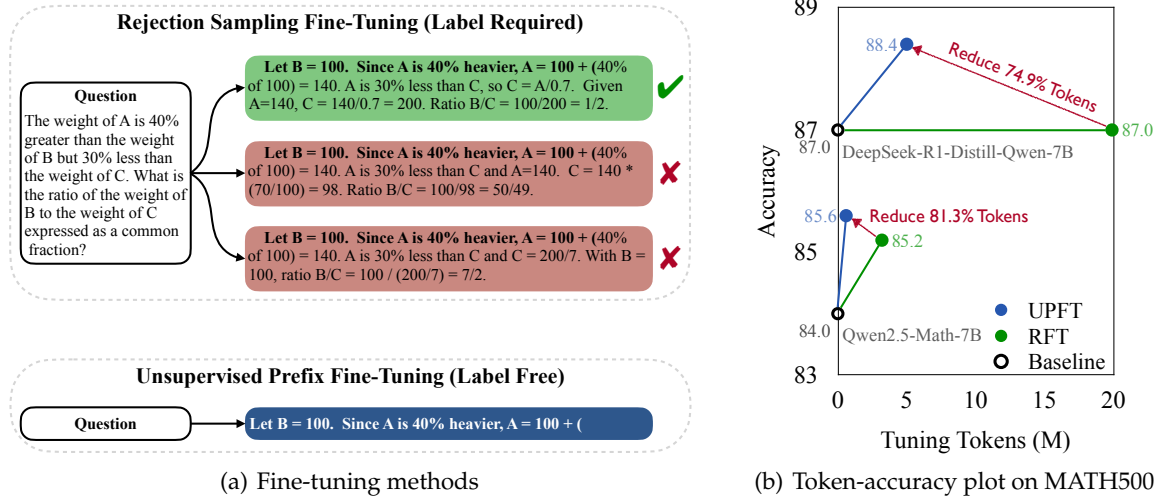


Figure 1: (a): Conventional Rejection Sampling Fine-Tuning (RFT) method (upper panel) involves generating multiple responses to a given question and then applying posterior filtering to discard trajectories that lead to incorrect answers. Finally, the correct trajectory is used for final training. In contrast, the proposed UPFT method (bottom panel) requires only prefix minimal initial tokens of a single generated sample, eliminating the need for labeled data or rejection sampling. (b): Our proposed UPFT matches the performance of supervised RFT, while reduces tuning cost by 75+%.

Abstract

Improving the reasoning capabilities of large language models (LLMs) typically requires supervised fine-tuning with labeled data or computationally expensive sampling. We introduce Unsupervised Prefix Fine-Tuning (UPFT), which leverages the observation of Prefix Self-Consistency – the shared initial reasoning steps across diverse solution trajectories – to enhance LLM reasoning efficiency. By training exclusively on the initial prefix substrings (as few as 8 tokens), UPFT removes the need for labeled data or exhaustive sampling. Experiments on reasoning benchmarks show that UPFT matches the performance of supervised methods such as Rejection Sampling Fine-Tuning, while reducing training time by 75% and sampling cost by 99%. Further analysis reveals that errors tend to appear in later stages of the reasoning process and that prefix-based training preserves the model’s structural knowledge. This work demonstrates how minimal unsupervised fine-tuning can unlock substantial reasoning gains in LLMs, offering a scalable and resource-efficient alternative to conventional approaches.

*Equal Contribution. The work was done when Ke Ji, Zhiwei He, Xingyu Chen, Xiaoyuan Liu, and Zhijie Wang were interning at Tencent AI Lab.

†Correspondence to: Zhaopeng Tu <zptu@tencent.com> and Benyou Wang <wangbenyou@cuhk.edu.cn>.

1 Introduction

Large language models (LLMs) have demonstrated remarkable performance across a wide range of natural language understanding and generation tasks, primarily due to large-scale pre-training and subsequent instruction fine-tuning on high-quality datasets (Longpre et al., 2023; Touvron et al., 2023). Despite these successes, enabling LLMs to exhibit systematic reasoning capabilities remains a challenging endeavor. In multiple domains—from mathematical problem solving to logical and commonsense reasoning—models often rely on large amounts of human-annotated data or extensive sampling-and-filtering pipelines to achieve high accuracy.

Recent inquiry has introduced approaches such as Rejection Sampling Fine-Tuning (RFT; Yuan et al. 2023) and Self-Taught Reasoner (STaR; Zelikman et al. 2022) to leverage model-generated solutions for iterative **self-improvement**, often requiring multiple candidate responses and subsequent filtering or verification steps. While these methods can yield impressive gains, they are time-consuming, resource-intensive, and assume ready access to correct targets or verification mechanisms — particularly challenging when no reliable ground-truth is available.

In this paper, we propose an unsupervised fine-tuning method that requires only a single pass of model-generated responses per question, coupled with prefix-based fine-tuning. Our key insight is that different solution paths often share a common initial reasoning phase, which we call “Prefix Self-Consistency”. By fine-tuning on these minimal prefixes (as few as 8 tokens), we effectively guide the model’s inherent reasoning structures toward more systematic solution strategies while avoiding the complexity and cost of large-scale or iterative filtering. Moreover, we preserve the model’s overall problem-solving format through a small proportion of full-token fine-tuning experiments, ensuring the model does not lose its length generalization and instruction-following abilities.

We conduct comprehensive experiments across four training corpora and evaluate our method on four widely used reasoning benchmarks. UPFT demonstrates exceptional data efficiency and flexibility, outperforming conventional full-token fine-tuning in an unsupervised sampling setting. Furthermore, UPFT achieves performance competitive with the widely used RFT method, which relies on labeled verification or large-scale rejection sampling, while requiring only 25% of the training time and 1% of the sampling time. Our approach can be easily adapted to various datasets, tasks, and LLM architectures, highlighting its flexibility and universality for developing robust problem-solving capabilities in large language models.

Our main contributions are as follows:

1. We identify Prefix Self-Consistency as a critical phenomenon, showing that early reasoning steps are highly consistent across trajectories, enabling efficient self-improvement learning.
2. We propose an unsupervised fine-tuning method UPFT that leverages only prefix substrings (i.e., minimal initial tokens of model-generated responses), eliminating the need for labeled data or rejection sampling.
3. We conduct comprehensive empirical validation to demonstrate that UPFT exhibits exceptional data efficiency and versatility, outperforming vanilla full-token fine-tuning in unsupervised settings and achieving competitive performance with RFT while drastically reducing sampling and tuning overhead.

2 Prefix Self-Consistency

A central observation of this work is that different solution trajectories for the same question often share a common *initial* reasoning phase, even if their later steps diverge substantially. We term this phenomenon **prefix self-consistency**. In other words, when an LLM generates multiple solutions for a single question, the opening tokens – typically restatements of the problem or the setup of its initial logical steps – tend to be highly consistent across these separate responses. In this section, we identify two characteristics of reasoning prefixes to demonstrate the existence and plausibility of

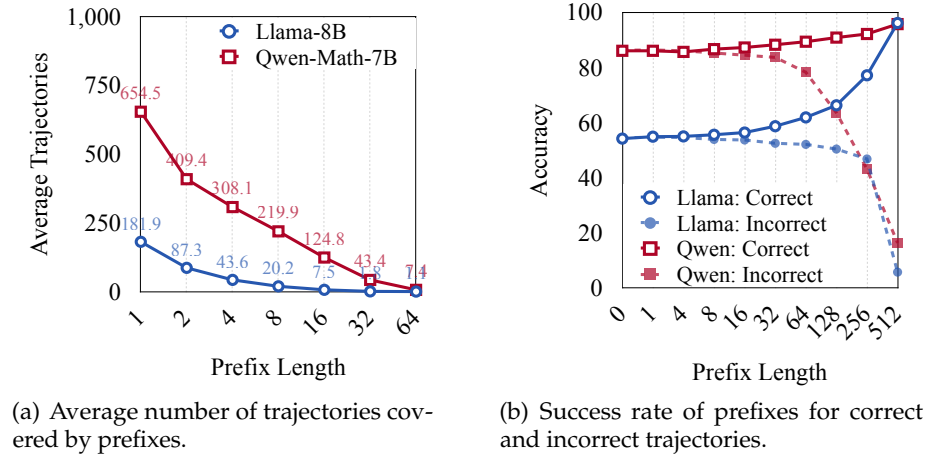


Figure 2: An empirical investigation of prefix self-consistency. We investigate (a) the average number of trajectories covered by prefixes at different lengths, and (b) the success rate of 32 rollout samplings from prefixes for both correct and incorrect trajectories.

prefix self-consistency. We report results on 500 questions randomly sampled from the PRM training dataset.

Early reasoning steps are highly consistent across reasoning trajectories. We sampled 1,000 trajectories for each question instance (using a temperature of 0.6) and calculated the average trajectories per prefix of different lengths, as shown in Figure 2(a). The results confirm that reasoning processes exhibit strong prefix self-consistency, with a remarkably high degree of prefix overlap among multiple trajectories for both models. As we increase t (i.e., consider longer prefixes), the average number of samples per prefix pattern decreases, yet both models maintain consistent patterns well beyond the very first tokens. Notably, the math-specialized Qwen-Math-7B-Instruct preserves shared prefixes more consistently than the general-purpose Llama-3.1-8B-Instruct, as evidenced by the higher values in all columns. This suggests that Qwen-Math-7B-Instruct’s generation process is more tightly anchored in its initial reasoning steps, whereas Llama-3.1-8B-Instruct introduces more prefix variability.

Errors predominantly occur in later reasoning steps. To empirically validate the plausibility of the reasoning prefix, we conducted 32 rollout samplings for each token in both a correct and an incorrect trajectory for each question. Figure 2(b) illustrates the success rate of these rollout samplings across both correct and incorrect trajectories. Two key observations can be made:

- *Incorrect trajectories diverge significantly from correct ones in later reasoning steps.* The rollout success rate for correct trajectories steadily rises as t grows for both models. For example, with Llama, the success rate starts at 54.2% at $t = 0$ and reaches 96.2% by $t = 512$. This trend is expected: sampling from the later tokens of a correct trajectory leverages more accurate contextual information, increasing the likelihood of staying on a correct path. In contrast, the rollout success rate for incorrect trajectories declines substantially as t increases, indicating that once an incorrect path is taken, recovering through rollout sampling becomes far less likely. This contrast highlights that errors tend to occur in the later stages of generation.
- *Initial reasoning prefixes exhibit self-consistency.* At the earliest token positions ($t \leq 16$), the rollout success rates for both correct and incorrect trajectories are strikingly similar. For Qwen-Math-7B-Instruct at $t = 4$, the success rates are 85.7% and 86.1% for correct and incorrect trajectories, respectively. This near-identical performance in early steps suggests that the initial tokens gener-

ated – whether they lead to a correct or incorrect final answer – are statistically indistinguishable in terms of leading to a correct result when used as a starting point for rollout sampling.

Overall, these results indicate that while mistakes in reasoning are more prone to appear and accumulate in later steps, the initial reasoning prefixes generated by LLMs exhibit a considerable degree of self-consistency. This consistency is reflected in the similar early-stage rollout success rates across both correct and incorrect trajectories. Consequently, enhancing the accuracy and robustness of these early reasoning steps may be a key factor in improving the overall reliability of complex reasoning tasks.

3 Methodology

Based on the aforementioned observation, we firstly introduce the conventional widely-used reject sampling strategy in Section 3.1, we then provide a Bayesian explanation and identify the coverage and accuracy of **Prefix Self-Consistency** in Section 3.2. Finally, based on this observation, we propose our UPFT in Section 3.3.

3.1 Preliminary

A prevalent strategy for enhancing reasoning in Large Language Models (LLMs) involves Supervised Fine-Tuning (SFT) using demonstrations of correct reasoning processes, often in the form of chain-of-thought (CoT) trajectories. This approach typically entails generating multiple reasoning traces from a base model and subsequently selecting only those traces where the extracted final answer aligns with the ground truth. These selected, correct reasoning trajectories are then used to fine-tune the base model via SFT. Given the inherent resemblance of this selection process to reject sampling, this methodology is also entitled with *Rejected Fine-Tuning* (RFT).

Formally, consider a dataset $(x, y) \sim \mathcal{D}$, where each x represents an input and y is the corresponding ground-truth answer. For each input x , we first generate a set of reasoning traces $\{r^{(k)}\}_{k=1}^K$ from the base model $p(\cdot|x; \theta)$ parameterized by θ . Then, for each input x , we filter these generated traces and uniformly sample a correct reasoning trace r from the set of traces whose final answer matches the ground truth y :

$$\begin{aligned} r^{(k)} &\sim p(\cdot|x; \theta) \\ r &\in_R \mathbb{1}(r^{(k)}, y) \end{aligned}$$

Function $\mathbb{1}(r, y)$ denotes the rejection sampling, and it is a selection function that only picks the reasoning trace r whose final answer is y , and \in_R denotes that the final selected reasoning trace r is uniformly selected from $\{r^{(k)}\}_{k=1}^K$. Consequently, the overall RFT objective is to maximize the following objective:

$$\mathbb{E}_{(x,y) \sim \mathcal{D}, r^{(k)} \sim p(\cdot|x; \theta)} \log p(r|x) \quad (1)$$

where the ground-truth answer y internally exists in the correct reasoning trace r .

3.2 Modeling Coverage and Accuracy

We demonstrate that the previous RFT process can be naturally interpreted within a Bayesian framework. Consider objective in Equation (1), the prediction of an answer y given an input x . From a probabilistic perspective, we can decompose the conditional probability $p(y|x)$ by marginalizing

over all possible reasoning traces r that could lead to the answer, which is:

$$\log p(y|x) = \log \sum_r p(y, r|x) \quad (2)$$

$$= \log \sum_r p(y|r, x) p(r|x) \quad (3)$$

$$\geq \sum_r p(r|x) \log p(y|r, x) \quad (4)$$

$$= \mathbb{E}_{r \sim p(\cdot|x)} \log p(y|r, x) \quad (5)$$

where we leverage the total probability rule in Eq.2, Bayes's Theorem in Eq.3, and Jensen's inequality in Eq.4, which finally yields an expectation term in Eq.5. It reveals that the log-likelihood of predicting the correct answer $\log p(y|x)$ is lower-bounded by two factors:

- **Coverage:** The expectation $\mathbb{E}_{r \sim p(\cdot|x)}[\cdot]$ is with respect to the distribution of reasoning trace $p(r|x)$. The term $p(r|x)$ denotes a prior probability of reasoning trace r given input x , which denotes the prefix coverage of the entire reasoning trace space. This corresponds to the average trajectory covered by prefixes in Section 2 indicated by Figure 2(a).
- **Accuracy:** Term $p(y|r, x)$ within the expectation can be viewed as the likelihood of the answer y being correct given a specific reasoning trace r , which is the accuracy. This is also observed by our previous observation in Section 2 in Figure 2(b).

Intuitively, a higher $p(r|x)$ for relevant reasoning traces that cover a broad range of problem-solving approaches indicates a broader exploration of potential solution paths. Simultaneously, a larger $\log p(y|r, x)$ implies that given a reasoning trace r , the model is more likely to arrive at the correct answer y .

Therefore, maximizing this lower bound necessitates enhancing both the coverage of a wide range of effective reasoning traces and the accuracy of each trace in reaching the desired solution. In terms of RFT, it is a specific example of trading off these two key factors: *RFT maximizes the accuracy of the reasoning trace by reject sampling while it neglects the coverage of the reasoning trace by only selecting one reasoning trace.*

3.3 Unsupervised Prefix Fine-Tuning

Given the observed dynamic shift from coverage to accuracy in previous content, in this subsection, we are motivated to investigate the existence of an optimal prefix length that balances both objectives in Equation (5).

Prefix Coverage and Accuracy To incorporate the concept of prefix spans, we utilize the chain rule for conditional probabilities. We can decompose the probability of the full reasoning trace r given input x as:

$$\begin{aligned} p(r|x) &= p(r_{<t}, r_{\geq t}|x) \\ &= p(r_{<t}|x) p(r_{\geq t}|r_{<t}, x) \end{aligned}$$

where $r_{<t}$ denotes the prefix of r before time step t . Substituting this decomposition of $p(r|x)$ into the original lower bound from Equation (4), we get (see Appendix A for the complete proof):

$$\log p(y|x) \geq \mathbb{E}_{r_{<t} \sim p(\cdot|x)} [L(r_{<t}, x)] \quad (6)$$

where $L(r_{<t}, x)$ is:

$$\mathbb{E}_{r_{\geq t} \sim p(\cdot|r_{<t}, x)} [\log p(y|r_{<t}, r_{\geq t}, x)] \quad (7)$$

Equation (6) presents the original Jensen's inequality lower bound in terms of prefix spans of the reasoning trace. It also states the importance of coverage and accuracy of the prefix span $r_{<t}$:

Task template used for prefix tuning.	
[question]	Please provide the initial step towards resolving the question. This step may serve as a foundation but might not encompass the entire solution.

Figure 3: The task template used to learn from the prefix of the reasoning traces. [question] represents the question that needs to be answered.

- **Prefix Coverage:** The outer expectation in Equation (6) $\mathbb{E}_{r_{<t} \sim p(\cdot|x)}[\cdot]$ is with respect to the distribution of prefixes $p(r_{<t}|x)$. The term $p(r_{<t}|x)$ represents the prior probability of the model generating a prefix reasoning trace $r_{<t}$ given the input x . This denotes the coverage of prefix reasoning traces.
- **Prefix Accuracy:** For each prefix $r_{<t}$, the term $L(r_{<t}, x)$ in Equation (7) is the conditional lower bound given the prefix $r_{<t}$. It is the expected log-likelihood of the answer y given that the reasoning process starts with prefix $r_{<t}$, averaging over all possible suffixes $r_{\geq t}$ that can follow $r_{<t}$. This can be viewed as the expected accuracy when reasoning is conditioned to start with prefix $r_{<t}$.

We have mathematically derived a representation of Jensen’s inequality lower bound in terms of prefix spans for both prefix coverage and prefix accuracy. To achieve superior performance through prefix fine-tuning, a specific prefix length t is required to trade off both prefix coverage and prefix accuracy terms.

Sampling Only First Few Tokens Motivated by this observation, we develop our proposed UPFT, which is apt to maximize the coverage of the reasoning trace while maintaining a relatively high accuracy by only learning from the proper prefix of the reasoning trace. This strategy shares several advantages. First, *UPFT exhibits enhanced coverage of all possible correct reasoning traces*. As demonstrated in Section 2 and Section 3.2, the prefix of the reasoning traces represents a set of reasoning traces with the identical prefix, which increases the coverage term of Equation (5)’s lower bound. Second, *by solely decoding a certain length of the prefix, UPFT inherently enjoys improved computational efficiency*. In contrast, conventional Reasoning From Traces (RFT) methods necessitate rejecting sampling over entire reasoning traces. This approach incurs a significant inference burden, particularly in scenarios where generating valid full reasoning traces from the base model proves challenging.

Our strategy consists of the following steps:

- Given a training set $(x, y) \in \mathcal{D}$, we only decode a prefix span $r_{<t}$ of reasoning trace from the base model $r_{<t} \sim p(\cdot|x; \theta)$;
- We conduct the SFT learning on the prefix spans of the reasoning trace $r_{<t}$ with NLL objective;

Structure Tuning To avoid the catastrophic forgetting (Muennighoff et al., 2025) of reasoning structures brought by prefix tuning, we adopt a simple yet effective multi-task learning approach, to jointly conduct the prefix tuning with conventional SFT process task. We use a probability ratio $p\%$ to randomly split out a subset $\mathcal{D}_f \subset \mathcal{D}$ to generate the full reasoning trace r for each data pair $(x, y) \in \mathcal{D}_f$, without justifying the correctness of this reasoning trace r or reject sampling process. Consequently, the data of the whole trace SFT is totally unsupervised, which also avoids the inference overhead of reject sampling. The rest of dataset \mathcal{D}_p is then used for prefix decoding $r_{<t}$. Moreover, to clearly demonstrate the base model that the learning process is solely based on prefix spans, we utilize a specialized task template, as shown in Figure 3, for prefix tuning to enhance the model’s ability to learn effectively from prefix-specific patterns.

Hyperparam.	PRM-12K	OMI2-60K	LIMO	U-Hard
Optimizer		AdamW		
Warmup Ratio		0.03		
Learning Rate		1e-6	2e-6	1e-6
LR Schedule		constant with warmup		
Batch Size		1		
Gradient Step		8		
Max Length		4096		16384
# Epoch	2	1	3	1

Table 1: The hyperparameters used for our method on all training corpora.

In summary, the overall objective is a combination of UPFT and unsupervised SFT by maximizing the following objective:

$$\mathbb{E}_{\substack{x \sim \mathcal{D}_p \\ r_{<t} \sim p(\cdot|x;\theta)}} [\log p(r_{<t}|x;\theta)] + \mathbb{E}_{\substack{x \sim \mathcal{D}_f \\ r \sim p(\cdot|x;\theta)}} [\log p(r|x;\theta)] \quad (8)$$

Notably, compared to conventional RFT objective, our method does not require any rejection sampling process denoted by $\mathbb{1}(r^{(k)}, y)$ in Equation (1). Consequently, our method is naturally suitable for learning from unproven questions with extreme data efficiency.

4 Experiment

In this experiment, we compare UPFT with traditional supervised methods such as SFT and RFT in both supervised and unsupervised sampling settings. We also evaluate its scalability by varying the self-training corpora and backbone models.

4.1 Experiments Setup

Backbone LLMs For our experiments, we selected three representative open-source backbone models: the general-purpose Llama-3.1-8B-Instruct (Dubey et al., 2024), the math-specialized Qwen2.5-Math-7B-Instruct (Yang et al., 2024) optimized for mathematical tasks, and the long-reasoning DeepSeek-R1-Distill-Qwen-7B (Guo et al., 2025), distilled from DeepSeek-R1.

Fine-Tuning We used four datasets to generate self-training data for fine-tuning:

1. **PRM** (12K instances; Lightman et al., 2024): This dataset includes 4.5K MATH (Hendrycks et al., 2021) test problems, which are drawn from the PRM800K training set.
2. **OMI2** (600K instances; Toshniwal et al., 2024): This is a subset of 600K unique questions extracted from the OpenMathInstruct2 math-instruction tuning dataset.
3. **LIMO** (819 instances; Ye et al., 2025): A high-quality training dataset specifically focused on challenging problems.
4. **U-Hard** (100K questions): This dataset, introduced in this work, is designed to explore the potential of our method in an unsupervised setting.

U-Hard was curated through an extensive collection of questions from publicly available online sources. Adhering to the Omni-Math approach (Gao et al., 2024), we labeled the collected data by difficulty and subsequently filtered out lower-difficulty questions. This process resulted in a dataset composed exclusively of challenging questions, thereby ensuring U-Hard’s practical relevance to real-world scenarios.

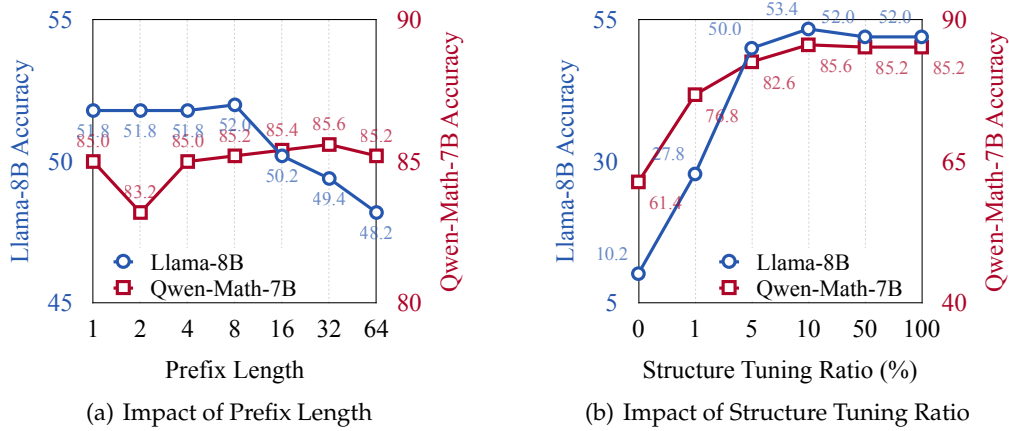


Figure 4: Impact of (a) prefix length and (b) structure tuning ratio on reasoning accuracy.

To ensure a fair comparison, UPFT employs the same hyperparameters as conventional SFT when fine-tuning models, as listed in Table 1. During the inference stage, we adopt a prompted zero-shot setup and use standard greedy decoding, wherein models are directed to answer each question using natural language instructions without any accompanying contextual demonstrations.

Benchmarks We evaluated the performance of our method and the baseline methods on four widely used benchmarks: **GSM8K** (Cobbe et al., 2021), **MATH500** (Hendrycks et al., 2021), **AIME24** (MAA Committees), and **GPQA Diamond** (Rein et al., 2023).

Tuning Scenarios We employed two experimental settings:

- **Unsupervised sampling:** we took one sample per question without any posterior filtering.
- **Supervised sampling:** Following RFT (Yuan et al., 2023), we sampled the responses for each question 16 times and used the ground truth answers to select a correct response.

4.2 Ablation Study

In this section, we evaluated the impact of two hyperparameters of UPFT on the PRM-12K dataset using Llama-3.1-8B-Instruct and Qwen2.5-Math-7B-Instruct in the unsupervised setting and report the results on the MATH500 test set.

Impact of Prefix Length Figure 4(a) illustrates how prefix length affects reasoning accuracy. Each model has its own optimal prefix length for peak performance: Llama-3.1-8B-Instruct achieves its highest average accuracy (52.0%) with an 8-token prefix, whereas Qwen2.5-Math-7B-Instruct remains stable over prefix lengths from 8 to 32 tokens. The math-specialized Qwen2.5-Math-7B-Instruct is less sensitive to prefix length, likely because it was trained on the high-quality large-scale math corpus, which reinforces its stability on the MATH500 test set. For subsequent experiments, we set Llama-3.1-8B-Instruct to 8 tokens and Qwen2.5-Math-7B-Instruct to 32 tokens. In the following experiments, we set the prefix length to 8 for Llama-3.1-8B-Instruct and 32 for Qwen2.5-Math-7B-Instruct. Because the responses from the long-reasoning DeepSeek-R1-Distill-Qwen-7B are generally more than four times longer than those of Qwen2.5-Math-7B-Instruct, we use a prefix length of 128 tokens for the former.

Impact of Structure Tuning Ratio Figure 4(b) plots the effect of the structure tuning ratio (p). The performance of both models generally improves with increasing p , peaking at $p = 10\%$. This result supports our hypothesis that even minimal structural supervision effectively guides models towards

Fine-Tuning			Testsets				
Method	Data	Avg. Length	GSM8K	MATH500	AIME2024	GPQA	Ave.
<i>Llama-3.1-8B-Instruct</i>			82.0	51.0	3.3	8.6	36.2
+ SFT	PRM	175.8	83.8	48.4	3.3	8.6	36.0
+ UPFT	(12K)	15.8	85.4	52.0	6.7	9.1	38.3
<i>Qwen2.5-Math-7B-Instruct</i>			95.2	84.0	16.7	9.6	51.4
+ SFT	PRM	300.1	95.8	83.4	13.3	9.1	50.4
+ UPFT	(12K)	51.4	95.5	85.6	20.0	9.6	52.6
+ SFT	OMI2	533.2	95.4	83.4	13.3	6.6	49.7
+ UPFT	(600K)	67.5	95.4	86.4	20.0	9.6	52.9
+ SFT	LIMO	491.8	95.8	84.2	20.0	7.6	51.9
+ UPFT	(0.8K)	77.8	95.6	85.8	20.0	8.6	52.5
+ SFT	U-Hard	393.3	95.5	83.4	16.7	9.6	51.3
+ UPFT	(100K)	68.2	96.0	85.6	26.6	9.6	54.5
<i>DeepSeek-R1-Distill-Qwen-7B</i>			88.6	87.0	40.0	13.1	57.2
+ SFT	LIMO	2029.5	89.7	87.0	40.0	12.1	57.2
+ UPFT	(0.8K)	757.7	92.0	89.4	43.3	17.7	60.6
+ SFT	U-Hard	3440.4	89.7	87.0	36.7	12.1	56.4
+ UPFT	(100K)	561.7	91.4	89.2	50.0	15.7	61.6

Table 2: Model performance under the **unsupervised sampling setting**, without any filtering based on the correctness of the extracted answer. “Length” denotes the average length of the tuning samples in each dataset. Models trained with SFT and the proposed UPFT generate responses of similar length.

generating complete responses. However, exceeding this ratio leads to a performance decrease. Consequently, we set p to 0.1 in UPFT for all datasets, except LIMO. For the smaller LIMO dataset, we increased p to 0.3 for structure tuning.

4.3 Unsupervised Fine-Tuning

Table 2 presents the results for the unsupervised sampling setting, where the data is not filtered based on the correctness of the extracted answer. We compare these outcomes to those obtained using conventional SFT (Ouyang et al., 2022), which fine-tunes models on training data without any self-improvement or test-time verification.

UPFT demonstrates superior performance compared to SFT in unsupervised fine-tuning. UPFT consistently outperforms conventional SFT across various self-training datasets and backbone models under the unsupervised sampling setting. For instance, using the U-Hard dataset with Qwen2.5-Math-7B-Instruct, UPFT achieves an average accuracy of 54.5% across all benchmarks, exceeding the 51.3% attained by conventional SFT. Likewise, with DeepSeek-R1-Distill-Qwen-7B and U-Hard, UPFT attains an average of 61.6%, whereas SFT achieves 56.4%. These results highlight UPFT’s effectiveness in leveraging unsupervised data to enhance reasoning, thereby reducing the reliance on labeled data. They also indicate the versatility and broad applicability of UPFT across different LLM architectures, supporting its claim of being a flexible, universal method for improving reasoning capabilities.

The benefits of UPFT are more pronounced on complex reasoning tasks. The performance gains of UPFT over conventional SFT are especially evident on more challenging benchmarks such as AIME2024 and GPQA. For example, on AIME2024 with the U-Hard dataset, UPFT achieves 26.6% accuracy with Qwen2.5-Math-7B-Instruct, a notable improvement over the 16.7% achieved by con-

Method	#Tokens		Testsets				
	Sampling	Tuning	GSM8K	MATH500	AIME2024	GPQA	Ave.
<i>Llama-3.1-8B-Instruct</i>			82.0	51.0	3.3	8.6	36.2
+ RFT		2.3M	86.0	52.0	6.7	9.1	38.5
+ V-STaR	36.9M	6.8M	85.4	52.6	6.7	8.6	38.3
+ UPFT (Ours)	0.2M		85.4	52.0	6.7	9.1	38.3
+ Lable Filter	36.9M	0.2M	85.8	53.4	6.7	9.1	38.8
<i>Qwen2.5-Math-7B-Instruct</i>			95.2	84.0	16.7	9.6	51.4
+ RFT		3.2M	95.7	85.2	20.0	9.6	52.6
+ V-STaR	51.7M	9.6M	96.0	85.4	20.0	10.1	52.9
+ UPFT (Ours)	0.6M		95.5	85.6	20.0	9.6	52.6
+ Lable Filter	51.7M	0.6M	96.0	85.6	20.0	10.1	52.9
<i>DeepSeek-R1-Distill-Qwen-7B</i>			88.6	87.0	40.0	13.1	57.2
+ RFT	318.0M	19.9M	90.7	87.0	40.0	11.1	57.2
+ UPFT (Ours)	5.0M		91.9	88.4	40.0	14.6	58.7
+ Lable Filter	318.0M	4.5M	92.3	89.2	40.0	13.6	58.8

Table 3: Model performance under the **supervised sampling setting** on the PRM-12K dataset, with filtering 16 sampled solutions based on the correct extract answer. “#Tokens” denote the number of tokens spent in each phase. Compared to RFT, V-STaR requires two additional samples to tune the verifier with DPO.

ventional SFT. A similar trend is observed with DeepSeek-R1-Distill-Qwen-7B on AIME2024, where UPFT reaches 50.0% compared with SFT’s 36.7%, showing UPFT’s capability to improve reasoning on difficult problems, aligning with our claim of enhancing reasoning capabilities effectively and efficiently.

The U-Hard dataset maximizes UPFT’s potential through difficulty-focused curation. Training with U-Hard yields the strongest performance improvements across models, particularly for complex benchmarks. Qwen2.5-Math-7B achieves 26.6% on AIME2024 with U-Hard - 10 points higher than with PRM data. This demonstrates that challenging questions provide richer signals for prefix-based self-improvement, as they demand more sophisticated initial reasoning setups. Using U-Hard, UPFT achieves the highest average accuracy with both Qwen2.5-Math-7B-Instruct and DeepSeek-R1-Distill-Qwen-7B, surpassing other datasets such as PRM-12K and OMI2-600K. These findings demonstrate that UPFT effectively leverages diverse and challenging questions to refine the model’s reasoning skills in an unsupervised manner, showcasing its data efficiency and versatility.

UPFT achieves dramatic efficiency gains through minimal token training. UPFT reduces training sequence length by 82.6-94.7% compared to SFT across datasets, with U-Hard samples averaging 68.2 tokens vs. SFT’s 393.3. This directly translates to 6.3-16.7x faster training iterations and reduced memory consumption. Notably, DeepSeek-R1-Distill-Qwen-7B attains better performance with UPFT’s 561-token samples than SFT’s 3,440-token sequences, proving that strategic prefix selection preserves critical learning signals. These efficiency characteristics validate UPFT’s practical value for resource-constrained applications while maintaining or exceeding baseline performance.

4.4 Supervised Fine-Tuning

In the supervised sampling setting, we compare our approach with two SFT variants that require labeled data in Table 3:

- RFT (Yuan et al., 2023) samples 16 candidate responses, then applies a label filter to identify a correct one.

- V-STaR (Hosseini et al., 2024) produces 16 responses for each tuning instance, trains a verifier on them for one iteration, and uses it at test time to select the answer from 4 completions.

UPFT achieves competitive performance with supervised methods while requiring significantly fewer tokens in both sampling and tuning. Across multiple model architectures, UPFT matches or exceeds the performance of supervised baselines such as RFT and V-STaR. On the Llama-3.1-8B-Instruct backbone, UPFT attains an average reasoning benchmark score of 38.3%, matching V-STaR (38.3%) and approaching RFT (38.5%). For Qwen2.5-Math-7B-Instruct, UPFT achieves identical average accuracy to RFT (52.6%), while using only 1.2% of the sampling tokens (0.6M vs. 51.7M). Most notably, UPFT enables DeepSeek-R1 to achieve 58.7% average accuracy, which is 1.5 points higher than RFT, while requiring $16\times$ fewer tuning tokens than RFT. This underscores how prefix-based fine-tuning effectively captures critical reasoning patterns without incurring the computational overhead of sampling 16 responses per question, as required by supervised baselines. Moreover, UPFT maintains its strong performance even when ground-truth answers are available. By focusing on the shared reasoning prefixes, it does not sacrifice accuracy. These results reinforce our core claim: prefix-based training captures the essential reasoning signals available in full trajectories, validating our Prefix Self-Consistency hypothesis. This hypothesis posits that the essential signals for effective reasoning are largely contained within the initial prefixes of successful solution trajectories, allowing for efficient learning and high performance.

UPFT offers seamless integration with label verification for enhanced accuracy. An additional benefit of UPFT is its inherent flexibility and compatibility with label verification techniques when such information is accessible. As demonstrated in the table’s last rows for each model backbone, when UPFT is enhanced with label filtering, it attains 58.8% average accuracy on DeepSeek-R1-Distill-Qwen-7B, surpassing all baselines while maintaining a $4.4\times$ tuning token advantage over RFT. For Qwen2.5-Math-7B-Instruct, we match V-STaR’s 52.9% peak performance without requiring compute-intensive verifier training or best-of-N inference. This dual capability demonstrates UPFT’s unique adaptability - it functions as a standalone unsupervised learner while seamlessly integrating supervision when available, making it practical for real-world deployment across the supervision spectrum.

5 Related Work

Large language models (LLMs) have achieved significant progress in natural language processing tasks (Longpre et al., 2023; Touvron et al., 2023). However, enabling LLMs to perform complex reasoning remains challenging.

Recent work has demonstrated strong reasoning capabilities in LLMs (Yang et al., 2024; Dubey et al., 2024; Guo et al., 2025; OpenAI, 2023). One research direction varies the input query of LLMs to obtain consistent responses through prompting techniques, such as chain-of-thought (Wei et al., 2022) and tree-of-thought (Yao et al., 2023). Another line of research leverages the verification of output trajectories for LLMs, for instance, ensembling and reranking inference paths, verification (Cobbe et al., 2021; Uesato et al., 2022) and self-consistency (Wang et al., 2023b). Trainable approaches like ReST (Zelikman et al., 2022), Rejection Sampling Fine-Tuning (RFT) (Yuan et al., 2023), and Self-Taught Reasoner (STaR) (Zelikman et al., 2022) also exemplify this paradigm.

However, a significant limitation of many existing approaches, particularly when considering complex tasks like mathematical reasoning, lies in their explicit reliance on sampling complete correct reasoning traces from LLMs. Although more recent studies such as s1 (Muennighoff et al., 2025) and LIMO (Ye et al., 2025) show that a small number of data can unlock strong generalization in LLMs for reasoning tasks, these approaches require external strong LLMs, which naturally are limited by external LLMs’ ability and cannot broaden the knowledge boundary by itself. Consequently, a critical question arises: How can we design an efficient and scalable method to enhance LLM reasoning without external supervision?

Our approach answers this question by leveraging the latent reasoning structures LLMs acquire during pre-training on large corpora (Brown et al., 2020), which can be aligned using a small subset of high-quality data (Muennighoff et al., 2025; Ye et al., 2025). We aim to enhance reasoning capabilities without annotated data by exploiting these inherent structures. Specifically, we observe that reasoning trajectories sampled by LLMs for the same question often share common prefix substrings, indicating a locally correct supervision signal, which we term Prefix Self-Consistency. Errors typically occur in later steps of the trajectory (see Appendix B for details).

Unsupervised Learning Unsupervised learning encompasses diverse methodologies, including pseudo-labeling (Ye et al., 2020), pivot-based approaches (Pan et al., 2010), and adversarial networks (Ganin et al., 2016). Self-supervised learning has since emerged as a dominant paradigm, particularly for language models, leading to the success of pre-trained models like BERT (Kenton & Toutanova, 2019) and others (Han et al., 2021; Radford et al., 2019). More recently, self-improvement techniques, such as self-rewarding (Chen et al., 2024; Yuan et al., 2024), further advance unsupervised learning by enabling models to iteratively enhance performance without external supervision. Building upon self-improvement, AL (Ji et al., 2024) extends this paradigm to unsupervised domain adaptation. In contrast to these general approaches, this paper focuses on discovering self-supervised signals specifically for mathematical reasoning. We introduce UPFT, a simple and effective unsupervised post-training method requiring only questions and the LLM itself.

Self-Training and Self-Improvement. A family of methods, starting with STaR (Zelikman et al., 2022), reinforced self-training (Gulcehre et al., 2023), and rejection fine-tuning (Yuan et al., 2023), leverages the solutions generated by large language models (LLMs) to iteratively update and improve the models themselves. These techniques involve fine-tuning the model on the generated solutions that produce correct answers. ReST^{EM} (Singh et al., 2023) views this fine-tuning process as expectation-maximization-based reinforcement learning (RL) for a solution-generating agent. Wang et al. (2023a) propose using a contrastive loss to enhance the likelihood of correct solutions over incorrect ones. The discovery of successful solutions presents a significant exploration challenge. Luong et al. (2024) demonstrated that RL-based fine-tuning of an LLM is particularly difficult unless preceded by several steps of supervised fine-tuning. In An et al. (2023), a more powerful LLM was employed to edit the incorrect rationales generated by a smaller model, thereby providing positive data for its fine-tuning. However, Huang et al. (2023) argued that LLMs have limited capacity to correct their own reasoning flaws.

6 Conclusion

In this work, we presented an unsupervised fine-tuning method that enhances the reasoning capabilities of large language models using only prefix substrings as minimal guidance. Our approach leverages the inherent reasoning structures within pretrained models, exploiting the phenomenon of Prefix Self-Consistency where different reasoning trajectories share common prefixes. Extensive experiments demonstrated that our method outperforms traditional full-token fine-tuning and achieves performance comparable to supervised approaches like RFT, with significantly reduced training and inference times. This work highlights the potential of minimal unsupervised fine-tuning in improving the reasoning abilities of LLMs without relying on external supervision or extensive computational resources. Future work will explore the application of this method to other challenging tasks and investigate the theoretical underpinnings of Prefix Self-Consistency in more depth.

References

Shengnan An, Zexiong Ma, Zeqi Lin, Nanning Zheng, Jian-Guang Lou, and Weizhu Chen. Learning from mistakes makes llm better reasoner. *arXiv preprint arXiv:2310.20689*, 2023.

- Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.
- Zixiang Chen, Yihe Deng, Huizhuo Yuan, Kaixuan Ji, and Quanquan Gu. Self-play fine-tuning converts weak language models to strong language models. In *ICML*, 2024.
- Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.
- Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.
- Yaroslav Ganin, Evgeniya Ustinova, Hana Ajakan, Pascal Germain, Hugo Larochelle, François Laviolette, Mario March, and Victor Lempitsky. Domain-adversarial training of neural networks. *Journal of machine learning research*, 17(59):1–35, 2016.
- Bofei Gao, Feifan Song, Zhe Yang, Zefan Cai, Yibo Miao, Qingxiu Dong, Lei Li, Chenghao Ma, Liang Chen, Runxin Xu, et al. Omni-math: A universal olympiad level mathematic benchmark for large language models. *arXiv preprint arXiv:2410.07985*, 2024.
- Caglar Gulcehre, Tom Le Paine, Srivatsan Srinivasan, Ksenia Konyushkova, Lotte Weerts, Abhishek Sharma, Aditya Siddhant, Alex Ahern, Miaosen Wang, Chenjie Gu, et al. Reinforced self-training (rest) for language modeling. *arXiv preprint arXiv:2308.08998*, 2023.
- Daya Guo, Dejian Yang, Haowei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu, Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- Xu Han, Zhengyan Zhang, Ning Ding, Yuxian Gu, Xiao Liu, Yuqi Huo, Jiezhong Qiu, Yuan Yao, Ao Zhang, Liang Zhang, et al. Pre-trained models: Past, present and future. *AI Open*, 2:225–250, 2021.
- Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.
- Arian Hosseini, Xingdi Yuan, Nikolay Malkin, Aaron Courville, Alessandro Sordoni, and Rishabh Agarwal. V-star: Training verifiers for self-taught reasoners. *arXiv preprint arXiv:2402.06457*, 2024.
- Jie Huang, Xinyun Chen, Swaroop Mishra, Huaixiu Steven Zheng, Adams Wei Yu, Xinying Song, and Denny Zhou. Large language models cannot self-correct reasoning yet. *arXiv preprint arXiv:2310.01798*, 2023.
- Ke Ji, Junying Chen, Anningzhe Gao, Wenya Xie, Xiang Wan, and Benyou Wang. Llms could autonomously learn without external supervision. *arXiv preprint arXiv:2406.00606*, 2024.
- Jacob Devlin Ming-Wei Chang Kenton and Lee Kristina Toutanova. Bert: Pre-training of deep bidirectional transformers for language understanding. In *Proceedings of NAACL-HLT*, pp. 4171–4186, 2019.
- Hunter Lightman, Vineet Kosaraju, Yuri Burda, Harrison Edwards, Bowen Baker, Teddy Lee, Jan Leike, John Schulman, Ilya Sutskever, and Karl Cobbe. Let’s verify step by step. In *The Twelfth International Conference on Learning Representations*, 2024.
- Shayne Longpre, Le Hou, Tu Vu, Albert Webson, Hyung Won Chung, Yi Tay, Denny Zhou, Quoc V Le, Barret Zoph, Jason Wei, et al. The flan collection: Designing data and methods for effective instruction tuning. 2023.

Trung Quoc Luong, Xinbo Zhang, Zhanming Jie, Peng Sun, Xiaoran Jin, and Hang Li. ReFT: Reasoning with reinforced fine-tuning. *arXiv preprint arXiv:2401.08967*, 2024.

MAA Committees. Aime problems and solutions. https://artofproblemsolving.com/wiki/index.php/AIME_Problems_and_Solutions.

Niklas Muennighoff, Zitong Yang, Weijia Shi, Xiang Lisa Li, Li Fei-Fei, Hannaneh Hajishirzi, Luke Zettlemoyer, Percy Liang, Emmanuel Candès, and Tatsunori Hashimoto. s1: Simple test-time scaling. *arXiv preprint arXiv:2501.19393*, 2025.

OpenAI. Gpt-4 technical report, 2023.

Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to follow instructions with human feedback. *Advances in Neural Information Processing Systems*, 35:27730–27744, 2022.

Sinno Jialin Pan, Xiaochuan Ni, Jian-Tao Sun, Qiang Yang, and Zheng Chen. Cross-domain sentiment classification via spectral feature alignment. In *Proceedings of the 19th international conference on World wide web*, pp. 751–760, 2010.

Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, Ilya Sutskever, et al. Language models are unsupervised multitask learners. *OpenAI blog*, 1(8):9, 2019.

David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. *arXiv preprint arXiv:2311.12022*, 2023.

Avi Singh, John D. Co-Reyes, Rishabh Agarwal, Ankesh Anand, Piyush Patil, Xavier Garcia, Peter J. Liu, James Harrison, Jaehoon Lee, Kelvin Xu, Aaron Parisi, Abhishek Kumar, Alex Alemi, Alex Rizkowsky, Azade Nova, Ben Adlam, Bernd Bohnet, Gamaleldin Elsayed, Hanie Sedghi, Igor Mordatch, Isabelle Simpson, Izzeddin Gur, Jasper Snoek, Jeffrey Pennington, Jiri Hron, Kathleen Kenealy, Kevin Swersky, Kshiteej Mahajan, Laura Culp, Lechao Xiao, Maxwell L. Bileschi, Noah Constant, Roman Novak, Rosanne Liu, Tris Warkentin, Yundi Qian, Yamini Bansal, Ethan Dyer, Behnam Neyshabur, Jascha Sohl-Dickstein, and Noah Fiedel. Beyond human data: Scaling self-training for problem-solving with language models. *arXiv preprint arXiv:2312.06585*, 2023.

Shubham Toshniwal, Wei Du, Ivan Moshkov, Branislav Kisacanin, Alexan Ayrapetyan, and Igor Gitman. Openmathinstruct-2: Accelerating ai for math with massive open-source instruction data. *arXiv preprint arXiv:2410.01560*, 2024.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023.

Jonathan Uesato, Nate Kushman, Ramana Kumar, Francis Song, Noah Siegel, Lisa Wang, Antonia Creswell, Geoffrey Irving, and Irina Higgins. Solving math word problems with process- and outcome-based feedback, 2022.

Peiyi Wang, Lei Li, Liang Chen, Feifan Song, Binghuai Lin, Yunbo Cao, Tianyu Liu, and Zhifang Sui. Making large language models better reasoners with alignment. *arXiv preprint arXiv:2309.02144*, 2023a.

Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc V Le, Ed H. Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. Self-consistency improves chain of thought reasoning in language models. In *The Eleventh International Conference on Learning Representations*, 2023b. URL <https://openreview.net/forum?id=1PL1NIMMrw>.

Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Fei Xia, Ed Chi, Quoc V Le, Denny Zhou, et al. Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems*, 35:24824–24837, 2022.

An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

Shunyu Yao, Dian Yu, Jeffrey Zhao, Izhak Shafran, Thomas L. Griffiths, Yuan Cao, and Karthik Narasimhan. Tree of thoughts: Deliberate problem solving with large language models, 2023.

Hai Ye, Qingyu Tan, Ruidan He, Juntao Li, Hwee Tou Ng, and Lidong Bing. Feature adaptation of pre-trained language models across languages and domains with robust self-training. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, pp. 7386–7399, 2020.

Yixin Ye, Zhen Huang, Yang Xiao, Ethan Chern, Shijie Xia, and Pengfei Liu. Limo: Less is more for reasoning. *arXiv preprint arXiv:2502.03387*, 2025.

Weizhe Yuan, Richard Yuanzhe Pang, Kyunghyun Cho, Sainbayar Sukhbaatar, Jing Xu, and Jason Weston. Self-rewarding language models. *arXiv preprint arXiv:2401.10020*, 2024.

Zheng Yuan, Hongyi Yuan, Chengpeng Li, Guanting Dong, Keming Lu, Chuanqi Tan, Chang Zhou, and Jingren Zhou. Scaling relationship on learning mathematical reasoning with large language models. *arXiv preprint arXiv:2308.01825*, 2023.

Eric Zelikman, Yuhuai Wu, Jesse Mu, and Noah D. Goodman. Star: Bootstrapping reasoning with reasoning. *Neural Information Processing Systems (NeurIPS)*, 2022.

A Proof of Prefix Lower bound

Following previous Equation (4), we have:

$$\begin{aligned}
 \log p(y|x) &\geq \sum_r p(r|x) \log p(y|r, x) \\
 &= \sum_r p(r_{<t}|x) p(r_{\geq t}|r_{<t}, x) \log p(y|r_{<t}, r_{\geq t}, x) \\
 &= \sum_{r_{<t}} \left[p(r_{<t}|x) \sum_{r_{\geq t}} p(r_{\geq t}|r_{<t}, x) \log p(y|r_{<t}, r_{\geq t}, x) \right] \\
 &= \sum_{r_{<t}} p(r_{<t}|x) \left[\sum_{r_{\geq t}} p(r_{\geq t}|r_{<t}, x) \log p(y|r, x) \right] \\
 &= \sum_{r_{<t}} p(r_{<t}|x) L(r_{<t}, x) \\
 &= \mathbb{E}_{r_{<t} \sim p(\cdot|x)} [L(r_{<t}, x)]
 \end{aligned} \tag{9}$$

where,

$$\begin{aligned}
 L(r_{<t}, x) &= \sum_{r_{\geq t}} p(r_{\geq t}|r_{<t}, x) \log p(y|r_{<t}, r_{\geq t}, x) \\
 &= \mathbb{E}_{r_{\geq t} \sim p(\cdot|r_{<t}, x)} [\log p(y|r_{<t}, r_{\geq t}, x)]
 \end{aligned} \tag{10}$$

B Case study of Prefix Data

We provide a case study of prefix data shown in Figure 5.

Case Study

Question: Melinda has three empty boxes and 12 textbooks, three of which are mathematics textbooks. One box will hold any three of her textbooks, one will hold any four of her textbooks, and one will hold any five of her textbooks. If Melinda packs her textbooks into these boxes in random order, the probability that all three mathematics textbooks end up in the same box can be written as $\frac{m}{n}$, where m and n are relatively prime positive integers. Find $m + n$.

The first 32 words (prefix substring) of all answers:

A1: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A2: To determine the probability that all three mathematics textbooks end up in the same box, we need to follow these steps: 1. **Calculate the total number of ways to pack the textbooks

A3: To determine the probability that all three mathematics textbooks end up in the same box, we start by calculating the total number of ways to distribute the 12 textbooks into the three

A4: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A5: To determine the probability that all three mathematics textbooks end up in the same box, we start by calculating the total number of ways to distribute the 12 textbooks into the three

A6: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A7: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A8: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A9: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A10: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A11: To determine the probability that all three mathematics textbooks end up in the same box, we start by calculating the total number of ways to distribute the 12 textbooks into the three

A12: To determine the probability that all three mathematics textbooks end up in the same box, we need to follow these steps: 1. **Calculate the total number of ways to distribute the 12

A13: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A14: To determine the probability that all three mathematics textbooks end up in the same box, we need to follow these steps: 1. **Calculate the total number of ways to pack the textbooks

A15: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

A16: To determine the probability that all three mathematics textbooks end up in the same box, we need to consider the total number of ways to distribute the 12 textbooks into the three

Figure 5: With the temperature set to 0.7, we sample 16 times based on Qwen2.5-Math-7B-Instruct for the given question, where A1-A16 represents the corresponding output results.